

Closed-View Auditing for Stateful Anonymity: Odometers, Alarms, and Machine-Checkable CPPCs

Unified archive note v0.08 (2026-02-28; archive v125)

Abstract

Deploying anonymous overlays and anonymous DHTs requires monitoring, debugging, and drift detection. But monitoring itself creates a control-plane leakage channel: logs, meters, and alerts reveal hidden state and can enable sequential identification. We propose a practical deployment pattern: a *measurement firewall* that keeps rich monitoring internal (closed view), and releases only budgeted, sanitized summaries governed by a privacy accountant (odometer/filter). We show that publicly publishing an alert bit (or trigger time/window) incurs an unavoidable privacy “oracle tax”: achieving low false-alarm and high power forces a minimum spend. Finally we propose CPPCs (Control-Plane Privacy Cards): short, machine-checkable manifests describing which witness channels are present, what accounting model is used, and what budget is allocated, analogous in spirit to model cards/datasheets and SBOM workflows (see ABOM/OINL, Anonymity C), but focused on anonymity control planes.[46]

Compilation note. This paper is brief-by-design. A companion addendum collects extended multi-stream monitoring notes and CPPC schema excerpts (State 3A).[47]

Scope. We isolate and formalize *Closed-View Auditing for Stateful Anonymity*: as a reusable contract surface in the anonymity / Anonymous-DHT series.

Units. Budgets are published in bits (use $\log_2$ and 2^ϵ); results stated in nats can be converted by dividing by $\ln 2$ (see Synthesis 8).[50]

Imports. We import the state-dependent anonymity viewpoint (amplification bounds, odometers/-filters as accountants) from AnonDHT State 1.[28] Stop-time padding and schedule compilation are treated as black-box mechanisms elsewhere in the series and appear here only as examples of witness channels CPPCs must disclose.[27, 26] CPPCs are designed to align with the series contract-object and manifest stack (Synthesis 0; ABOM/OINL, Anonymity C).[45, 46] Trace-interface naming (exposure normal forms, ENF-ID) and threat-window declarations (TW-ID) are imported from Synthesis 3/4.[48, 49]

Exports.

Claim fields (non-negotiable). A CPPC must record which transcript projection it is constraining (`exposure_nf_id`) and which repetition/threat window the bound is meant to cover (`tw_id`). For stateful monitoring surfaces, CPPCs should also carry a digest-bound `state_decl_id` committing to the declared state surfaces and accounting scope (see State 1 and the worked example, Synthesis 17).

Figure omitted in source-only archive.
(Plot/diagram and scripts available on request.)

Figure 1: Measurement firewall pattern. Rich telemetry stays internal; releases are budgeted and coarse.

Figure omitted in source-only archive.
(Plot/diagram and scripts available on request.)

Figure 2: Closed-view composition stack.

These ids are claim fields, not prose: they are carried by the contract object, by receipts/certificates, and by the ABOM that commits the supporting bundle.

We export (i) the *measurement firewall* pattern for keeping rich monitoring closed-view while releasing only budgeted summaries, (ii) an explicit “alert-oracle tax” lower bound (Cor. 1) quantifying the unavoidable spend of publishing alarms, and (iii) a CPPC (Control-Plane Privacy Card) manifest schema for machine-checkable disclosure of witness channels, accountants, and allocated budgets, treated as a first-class contract object in the archive interface (Synthesis 0).[45].

Artifact policy. This archive is source-only; supporting artifacts (plots, scripts, datasets, and tooling) are available on request.[44]

1 Introduction

Anonymity systems require monitoring: to detect congestion collapse, path bias, Sybil capture, or failures. Yet monitoring is itself a control-plane channel. Publishing fine-grained telemetry, alarms, or “health” beacons can enable sequential identification of hidden state.

Closed view. We propose a deployment pattern: keep rich telemetry and detectors *inside* a measurement firewall, and release only budgeted, sanitized summaries. The core idea is to treat monitoring as an interactive mechanism that must be accounted for.

Non-redundancy. AnonDHT State 1 develops state-dependent anonymity games and amplification bounds [28]. Here we focus on the deployment interface: how to monitor without creating a public witness stream, and how to document/attest the control-plane privacy contract via CPPC manifests.

2 Closed-view architecture

Composition stack. A useful mental model is a composition stack: raw telemetry feeds detectors; detectors feed internal dashboards; and only selected, budgeted releases cross the firewall.

Why “closed view” is technically meaningful. If the public sees only a coarsened or sub-sampled view of internal events, then the observer channel can contract distinguishability. In DP analyses, this corresponds to privacy amplification by subsampling/coarsening [29]. The firewall enforces this by policy: detectors may look at everything internally, but releases are constrained.

Figure omitted in source-only archive.
 (Plot/diagram and scripts available on request.)

Figure 3: Illustrative alarm tax curve: tighter (α, β) implies larger minimum ε .

Accountant placement. An odometer/filter is placed at the firewall boundary. It tracks realized privacy loss (odometer) or halts publication (filter) when a budget is exceeded [8]. We emphasize adaptive-parameter regimes: operators often increase monitoring or tighten thresholds in response to events.

3 Accounting, alarms, and the alert-oracle tax

Public alerts are witnesses. An “alarm bit” is a hypothesis test output; publishing the trigger time/window turns monitoring into a sequential test channel. This issue is explicit in the DP sequential detection literature: releasing a detector statistic and its stopping time can compromise privacy, motivating DP variants that noise both statistics and thresholds [32].

AnonDHT State 1 develops the general “stability $\Rightarrow$ distinguishing advantage” bridge for state-dependent anonymity channels; here we apply the same logic to the simplest possible public witness: an alarm bit.

3.1 A minimal tax statement

The relationship between DP-style stability and hypothesis testing is classical; see, e.g., max-divergence/DP as a bound on likelihood ratios and testing power [1, 3].

Corollary 1 (Alert-oracle tax). *Let $B \in \{0, 1\}$ be a public alarm bit released by a detector, and let $\sigma \in \{0, 1\}$ index two adjacent/control-plane states (e.g., “good” vs “bad”). If the detector has false-alarm probability $\Pr[B=1 \mid \sigma=0] \leq \alpha$ and miss probability $\Pr[B=0 \mid \sigma=1] \leq \beta$, then any (ε, δ) stability guarantee for the alarm channel must satisfy*

$$1 - \beta \leq 2^\varepsilon \alpha + \delta.$$

In particular, if $1 - \beta > \delta$, then $\varepsilon \geq \log_2((1 - \beta - \delta)/\alpha)$.

Proof. Apply the (ε, δ) inequality to the event $E = \{B = 1\}$: $\Pr[B=1 \mid \sigma=1] \leq 2^\varepsilon \Pr[B=1 \mid \sigma=0] + \delta$. Use $\Pr[B=1 \mid \sigma=1] \geq 1 - \beta$ and $\Pr[B=1 \mid \sigma=0] \leq \alpha$. $\square$

Figure 3 provides an illustration.

Concrete calibration (brief). Table 1 gives a numeric sense of Corollary 1; tight public alarms imply nontrivial ε spending.

3.2 Mitigation patterns for public reporting

The tax in Corollary 1 is a design constraint: if you insist on publishing *high-quality* alerts, you are committing to spending a nontrivial privacy/stability budget on the alert channel. Conversely, if you must keep the public budget small, you must deliberately degrade the alert channel (larger α and/or β).

A few practical patterns are simple and composable:

Miss prob. β	False-alarm α	Lower bound $\varepsilon \gtrsim \log_2((1 - \beta)/\alpha)$
0.10	0.05	2.89
0.10	0.01	4.50
0.10	0.001	6.80
0.05	0.05	2.94
0.05	0.01	4.55
0.05	0.001	6.86
0.01	0.05	2.99
0.01	0.01	4.60
0.01	0.001	6.90

Table 1: Calibration for Corollary 1: tight alarms cost nontrivial ε . This table is illustrative and ignores small- δ and composition details.

1. **Fixed schedule (no data-dependent timing).** Release summaries at predetermined times; do not publish “as soon as” a detector fires.
2. **Coarsen and delay.** Map the trigger time/window T to a coarse bucket $B = f(T)$ and release only after a fixed delay. Post-processing cannot worsen stability/privacy: if an internal report is (ε, δ) -DP then so is any $f(\cdot)$ of it.
3. **Delayed aggregates under continual observation.** When possible, replace per-event public alarms with aggregate releases (counts/rates) on a fixed cadence, using continual-observation DP mechanisms (binary-tree / hybrid counters). This trades a sharp stopping-time oracle for a smoothed stream whose disclosure properties are well understood [4, 5, 6, 7].
4. **Output perturbation.** If an alert must be public, add symmetric randomized-response noise to the published bit; this raises false alarms by design and can enforce a target (α, β) curve under a fixed ε budget [42].
5. **Commit-then-reveal (fixed time).** When operators want accountability without a real-time oracle, publish a daily/weekly commitment to the internal alert transcript (e.g., a hash or Pedersen-style commitment) on a fixed schedule, and reveal the plaintext transcript only after a policy delay. The commitment anchors integrity (prevents backdating) while the delayed reveal intentionally degrades the public channel by widening the effective alert window, aligning with Corollary 1 [43].

These patterns are intentionally generic: they can be layered on top of any closed-view detector, and they make the “alert quality vs. public budget” tradeoff explicit.

Lemma 1 (Randomized-response degradation of a public alarm). *Let $B \in \{0, 1\}$ be an internal alarm bit with false-alarm rate $\alpha_0 = \Pr[B=1 \mid \sigma=0]$ and miss rate $\beta_0 = \Pr[B=0 \mid \sigma=1]$. Publish $\tilde{B} := B \oplus N$ where $N \sim \text{Bernoulli}(q)$ is independent symmetric noise. Then the public false-alarm and miss rates are*

$$\alpha = (1 - q)\alpha_0 + q(1 - \alpha_0), \quad \beta = (1 - q)\beta_0 + q(1 - \beta_0).$$

In particular, increasing q intentionally degrades alert quality (driving (α, β) toward $(1/2, 1/2)$), which can be used to satisfy a fixed public budget implied by Corollary 1.

Figure omitted in source-only archive.
(Plot/diagram and scripts available on request.)

Figure 4: Example of an anytime-valid monitoring boundary for adaptive/streaming settings.

Figure omitted in source-only archive.
(Plot/diagram and scripts available on request.)

Figure 5: A “compiler” view of CPPCs: selection + releases + monitoring are compiled into a single budgeted interface.

3.3 Anytime-valid monitoring and many alarms

Closed-view detectors can be anytime-valid (confidence sequences, e-processes) [13, 14]. The default Ville/Markov rejection threshold $1/\alpha$ for e-processes is often conservative; recent work gives sharper thresholds under mild, empirically checkable shape assumptions on the e-value distribution [17]. But if many detectors exist, naive “raise any alarm” logic incurs multiple-testing inflation. E-values provide dependence-robust FDR control (e-BH) [15]. When alarms are raised at global stopping times, one should use stopped e-BH variants [16].

4 CPPC manifests

We propose **Control-Plane Privacy Cards (CPPCs)**: short, machine-checkable manifests describing control-plane witness channels and their budgets. A CPPC is not a proof; it is a contract that enables third-party auditing and comparison.

What a CPPC declares. At minimum: (i) which control-plane signals are released (weights, beacons, alerts), (ii) cadence and granularity, (iii) accounting model (e.g., max-divergence DP, RDP, or tradeoff-curve notions such as f -DP/GDP [36]) and composition regime, (iv) refresh/expiration policies, (v) how monitoring/alarms are handled (closed view vs public), (vi) the *unit of privacy* / adjacency notion and any interpretability guidance for the chosen budget representation (e.g., as emphasized by NIST’s guidance for evaluating DP guarantees [31]).

Publication and integrity. In practice, CPPCs should be signed and published via an append-only transparency mechanism (Merkle-tree log), analogous to Certificate Transparency [40]. This is also in the spirit of “documentation artifacts” such as model cards and datasheets [19, 18], and can reuse supply-chain primitives such as SBOM formats (e.g., SPDX or CycloneDX) [24, 23].

Verification sketch. A minimal deployment can be lightweight: (i) the operator signs a CPPC (and version bumps) and submits it to a public log; (ii) each public release channel (weights, beacons, alerts) includes a short log pointer (e.g., inclusion proof or log checkpoint); (iii) auditors verify the signature and log inclusion, then check *consistency* between CPPC declarations and observable releases: cadence/granularity match; budget beacons are monotone and compose according to the declared accountant; and any public alerting behavior matches the declared “closed view vs public” handling. This does not reveal raw telemetry, but it makes covert, out-of-contract control-plane channels harder to hide.

Figure omitted in source-only archive.
(Plot/diagram and scripts available on request.)

Figure 6: Example CPPC cards.

Auditor can verify from public artifacts	Typically not verifiable without trust/attestation
CPPC signature validity and log inclusion (inclusion/consistency proofs) [40, 41, 20]	Correctness of internal telemetry, sampling, and detector implementations
Release cadence and granularity match CPPC (timestamps, bucket sizes)	Whether additional closed-view channels exist beyond those declared
Budget beacon monotonicity and declared composition schedule (odometer/filter)	The operator’s internal “stop time” logic (e.g., whether early alarms were suppressed)
Public alert policy matches CPPC (fixed schedule vs data-dependent, use of randomized response)	Confidential incident details and on-call run-books
Artifacts referenced by the CPPC (SBOM-like metadata, versioning) are present and consistent [25, 22, 21]	Whether budgeted releases were computed from the claimed inputs (needs reproducibility or trusted execution)

Table 2: A minimal checklist for “budget-consistency” auditing. The goal is not to prove security from public data, but to make it difficult to silently change witness channels, cadences, or budgets without leaving detectable inconsistencies.

5 Related work

Accounting. Differential privacy provides the canonical stability language for releasing summaries [1]. Odometers and filters address adaptive-parameter composition [8], with practical RDP-oriented variants [10, 2]. Interactive deployments motivate concurrent composition results [9]. Privacy amplification by subsampling supports the “closed-view” intuition that partial/public views can be strictly safer [29]. Tradeoff-curve notions such as f -DP (and its Gaussian specialization, GDP) preserve the testing interpretation and enable sharp composition reasoning [36]. Recent work studies when fully-adaptive f -DP filters are valid and gives approximate GDP filters for practical regimes [37]. Natural (profile-level) privacy filters promise sharper accounting but are *not* always “free” under adaptive composition, even for common families of mechanisms [38]. “Accuracy-first” and ex-post privacy perspectives offer another interface for adaptively choosing privacy parameters while preserving post-processing immunity in the accounting model [39].

Always-valid monitoring and multiple testing. Confidence sequences and e-values support anytime-valid inference in streaming systems [13, 14]. E-values yield dependence-robust FDR control via e-BH [15]; stopped e-BH variants address global stopping-time alerts [16]. Recent work applies these tools to *audit* privacy guarantees from sampled outputs, including sequential auditors for (ϵ, δ) -DP and for f -DP [33, 34, 35].

Private telemetry systems. Prio and RAPPOR demonstrate scalable private aggregation and telemetry in production-like settings [12, 11].

Transparency and “cards.” Model cards and datasheets helped normalize lightweight, comparable documentation for models/datasets [19, 18]. Supply-chain work provides analogous machine-checkable manifests (SBOM) and provenance: SPDX and CycloneDX, SLSA (v1.0 is marked as “retired as of 2026), and in-toto [25, 22, 21]. Append-only transparency logs (Certificate Transparency) provide a simple publication primitive [40], while Sigstore aims to make signing/provenance usable at scale [20]. NIST’s Secure Software Development Framework (SSDF) is a recent baseline for secure development practice [30]. CPPCs adapt these documentation/provenance ideas specifically to anonymity control planes.

6 Takeaways

- **Monitoring is a witness stream:** alarms, trigger times, and public health beacons can act as sequential identification oracles and must be priced.
- **Prefer closed view:** keep rich telemetry and detectors inside a measurement firewall; release only coarse, precommitted summaries governed by an accountant.
- **Account adaptivity explicitly:** use odometers/filters (or f -DP/GDP tradeoff accounting) when thresholds, cadences, or detector libraries change over time.
- **Make the contract machine-checkable:** CPPCs and transparency logging make it harder to silently change witness channels, cadences, or budgets without leaving a detectable inconsistency.

References

- [1] Cynthia Dwork and Aaron Roth. The Algorithmic Foundations of Differential Privacy. *Foundations and Trends in Theoretical Computer Science*, 9(3–4):211–407, 2014.
- [2] Ilya Mironov. Rényi Differential Privacy. In *CSF*, 2017. [arXiv:1702.07476](#).
- [3] Borja Balle. Hypothesis Testing Interpretations and Rényi Differential Privacy. In *ICML (PMLR 108)*, 2020.
- [4] Cynthia Dwork, Moni Naor, Toniann Pitassi, and Guy Rothblum. Differential Privacy Under Continual Observation. In *STOC*, 2010.
- [5] T.-H. Hubert Chan, Elaine Shi, and Dawn Song. Private and Continual Release of Statistics. [ePrint:2010/076](#). (See also TCC/ICALP-era versions.)
- [6] Palak Jain, Sofya Raskhodnikova, Satchit Sivakumar, and Adam Smith. The Price of Differential Privacy under Continual Observation. In *ICML (PMLR 202)*, 2023.
- [7] Edith Cohen, Xin Lyu, Jelani Nelson, Tamás Sarlós, and Uri Stemmer. Lower Bounds for Differential Privacy Under Continual Observation and Online Threshold Queries. In *ICML (PMLR 247)*, 2024.
- [8] Ryan Rogers, Aaron Roth, Jonathan Ullman, and Salil Vadhan. Privacy Odometers and Filters: Pay-as-you-Go Composition. In *NeurIPS*, 2016. [arXiv:1605.08294](#).
- [9] Salil Vadhan and Tianhao Wang. Concurrent Composition of Differential Privacy. In *TCC*, 2021.
- [10] Mathias Lécuyer et al. Practical Privacy Filters and Odometers with Rényi Differential Privacy (and Applications to DP Deep Learning). [arXiv:2103.01379](#), 2021.

- [11] Úlfar Erlingsson, Vasyl Pihur, and Aleksandra Korolova. RAPPOR: Randomized Aggregatable Privacy-Preserving Ordinal Response. In *CCS*, 2014. [arXiv:1407.6981](#).
- [12] Henry Corrigan-Gibbs and Dan Boneh. Prio: Private, Robust, and Scalable Computation of Aggregate Statistics. In *NSDI*, 2017. [arXiv:1703.06255](#).
- [13] Steven R. Howard et al. Time-uniform, nonparametric, nonasymptotic confidence sequences. *Annals of Statistics*, 49(2), 2021. [arXiv:1810.08240](#).
- [14] Aaditya Ramdas. Hypothesis Testing with E-values. [arXiv:2410.23614](#), 2024.
- [15] Ruodu Wang. False discovery rate control with e-values. *JRSS B*, 84(3):822–852, 2022.
- [16] Hongjian Wang, Sanjit Dandapanthula, and Aaditya Ramdas. Anytime-valid FDR control with the stopped e-BH procedure. [arXiv:2502.08539](#), 2025.
- [17] Christopher Blier-Wong and Ruodu Wang. Improved thresholds for e-values. Preprint / early 2026 versions available online.
- [18] Timnit Gebru et al. Datasheets for Datasets. *Communications of the ACM*, 64(12), 2021. (also [arXiv:1803.09010](#)).
- [19] Margaret Mitchell et al. Model Cards for Model Reporting. In *FAT**, 2019. [arXiv:1810.03993](#).
- [20] Zachary Newman et al. Sigstore: Software Signing for Everybody. In *CCS*, 2022.
- [21] Santiago Torres-Arias et al. in-toto: Providing farm-to-table guarantees for bits and bytes. In *USENIX Security*, 2019.
- [22] SLSA. SLSA specification v1.0 (Status: Retired). <https://slsa.dev/spec/v1.0/>.
- [23] OWASP CycloneDX. Specification overview. <https://cyclonedx.org/specification/overview/>.
- [24] SPDX. SPDX specification overview. <https://spdx.github.io/spdx-spec/v3.0.1/>.
- [25] SPDX. SPDX Specification v3.0.1. <https://spdx.dev/wp-content/uploads/sites/31/2024/12/SPDX-3.0.1-1.pdf>.
- [26] Anonymous. *Scheduling and Compiler Techniques for Metadata-Hiding Overlay Lookups*. Series paper (published), 2026. Wiki: [\[\[2026.01.22 - Mathematics: Scheduling and Compiler Techniques for Metadata-Hiding Overlay Lookups\]\]](#).
- [27] Anonymous. *Stop-Time Padding Addendum (unique-only): Max-Mixture Separations and Tail-Sign Deadline Normal Forms*. Series note (published), 2026. Wiki: [\[\[2026.01.26 - Mathematics: Stop-Time Padding Addendum, Max-Mixture Separations and Tail-Sign Deadline Normal Forms\]\]](#).
- [28] Anonymous. Stateful Anonymity and Control-Plane Leakage: Selection Privacy, Amplification, and Accounting. Companion preprint in this archive (AnonDHT State 1).
- [29] Borja Balle and Gilles Barthe and Marco Gaboardi. Privacy Amplification by Subsampling: Tight Analyses via Couplings and Divergences. *Advances in Neural Information Processing Systems (NeurIPS)*, 2018. NeurIPS 2018; also available as [arXiv:1807.01647](#)
- [30] National Institute of Standards and Technology. Secure Software Development Framework (SSDF) Version 1.1: Recommendations for Mitigating the Risk of Software Vulnerabilities. NIST Special Publication 800-218, 2022. <https://csrc.nist.gov/pubs/sp/800/218/final>

- [31] Joseph P. Near and David Darais and Naomi Leffkowitz and Gary S. Howarth. Guidelines for Evaluating Differential Privacy Guarantees. NIST Special Publication 800-226, 2025. DOI: 10.6028/NIST.SP.800-226. Final publication (March 2025) with PDF available from the NIST landing page. <https://csrc.nist.gov/pubs/sp/800/226/final>
- [32] Liyan Xie and Ruizhi Zhang. Sequential Change Detection with Differential Privacy. 2025. DOI: 10.48550/arXiv.2509.02768. arXiv:2509.02768 <https://arxiv.org/abs/2509.02768>
- [33] Tomás González and Mateo Dulce-Rubio and Aaditya Ramdas and Mónica Ribero. Sequentially Auditing Differential Privacy. 2025. arXiv:2509.07055; also presented at NeurIPS 2025. <https://arxiv.org/abs/2509.07055>
- [34] Tim Kutta and Martin Dunsche and Yu Wei and Vassilis Zikas. Sequential Auditing for f -Differential Privacy. 2026. arXiv:2602.06518
- [35] Saeed Mahloujifar and Luca Melis and Kamalika Chaudhuri. Auditing f -Differential Privacy in One Run. Proceedings of the 42nd International Conference on Machine Learning (ICML), 267, 42615–42641, 2025. Also available as arXiv:2410.22235 and via OpenReview. <https://proceedings.mlr.press/v267/mahloujifar25a.html>
- [36] Jinshuo Dong and Aaron Roth and Weijie J. Su. Gaussian Differential Privacy. Journal of the Royal Statistical Society: Series B (Statistical Methodology), 84(1), 3–37, 2022. DOI: 10.1111/rssb.12455.
- [37] Long Tran and Antti Koskela and Ossi Räsä and Antti Honkela. f -Differential Privacy Filters: Validity and Approximate Solutions. 2026. DOI: 10.48550/arXiv.2602.06756. arXiv:2602.06756 <https://arxiv.org/abs/2602.06756>
- [38] Matthew Regehr and Bingshan Hu and Ethan Leeman and Pasin Manurangsi and Pierre Tholoni and Mathias Lécuyer. Natural Privacy Filters Are Not Always Free: A Characterization of Free Natural Filters. 2026. DOI: 10.48550/arXiv.2602.15815. arXiv:2602.15815 <https://arxiv.org/abs/2602.15815>
- [39] Ossi Räsä and Antti Koskela and Antti Honkela. Accuracy-First Rényi Differential Privacy and Post-Processing Immunity. 2025. DOI: 10.48550/arXiv.2509.22213. arXiv:2509.22213 <https://arxiv.org/abs/2509.22213>
- [40] Ben Laurie and Adam Langley and Emilia Kasper and Eran Messeri. Certificate Transparency Version 2.0. 2021. RFC 9162. <https://www.rfc-editor.org/rfc/rfc9162.html>
- [41] Alin Tomescu and Vivek Bhupatiraju and Dimitrios Papadopoulos and Charalampos Papamanthou and Nikos Triandopoulos and Srinivas Devadas. Transparency Logs via Append-Only Authenticated Dictionaries. Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security (CCS '19), 1299–1316, 2019. DOI: 10.1145/3319535.3345652. General-purpose transparency log construction and security formalization.
- [42] Stanley L. Warner. Randomized Response: A Survey Technique for Eliminating Evasive Answer Bias. Journal of the American Statistical Association, 60(309), 63–69, 1965. DOI: 10.1080/01621459.1965.10480775.
- [43] Torben P. Pedersen. Distributed Provers with Applications to Undeniable Signatures. Advances in Cryptology — EUROCRYPT '91, 547, 221–242, 1991. DOI: 10.1007/3-540-46416-6_20.
- [44] Anonymous. Artifact policy and supporting materials for the anonymity/AnonDHT series. 2026. See synthesis paper: Artifact policy (this archive). Artifacts, tooling, certificates, and scripts are available on request as a single archive.
- [45] Anonymous. Contract Objects for Anonymous-DHT Privacy Budgets and Claims. Series synthesis note (Synthesis 0), 2026. (This archive.)

- [46] Anonymous. ABOM/OINL: Audit BOM and Outcome Identification Noise Label for Anonymous DHT Deployments. Anonymity series Paper C, 2026. (This archive.)
- [47] Anonymous. Closed-View Auditing and CPPCs (Addendum): Multi-Stream Monitoring Notes and Schema Excerpts. AnonDHT State series Paper 3A, 2026. (This archive.)
- [48] Anonymous. Trace Interfaces for Anonymous DHTs: Observation Tiers, Committee-Contact Traces, and Exposure Normal Forms. Series synthesis note (Synthesis 3), 2026. (This archive.)
- [49] Anonymous. Threat Windows for Anonymous DHTs: What a Budget Covers and For How Long. Series synthesis note (Synthesis 4), 2026. (This archive.)
- [50] Anonymous. Notation and Minimal Model for the One-True-Archive. Series synthesis note (Synthesis 8), 2026. (This archive.)